

Technical Report: Advanced Validation, Overfitting Control & Hyperparameter Tuning

Prepared by: P.Chandan

Project: Artificial Intelligence & Machine Learning – Task 3

Dataset: California Housing Dataset

1. Executive Summary

Building upon the data preprocessing and multi-model baseline established in Task 2, this report details the transition of the California House Price Prediction system into an industry-grade framework. While prior milestones focused primarily on achieving baseline metric scores on a single train-test split, Task 3 shifts the engineering focus toward model reliability and operational generalization. Through the implementation of systematic overfitting diagnostics, 5-Fold Cross-Validation, and automated hyperparameter optimization using GridSearchCV, a highly stable and non-linear regression model was successfully built, verified, and justified for production-level reliability.

2. Problem Statement: The Challenge of Overfitting

In professional data science deployment, high accuracy on training data is a misleading metric if the model cannot adapt to uncorrupted, real-world data. Unconstrained non-linear estimators, such as a default Decision Tree Regressor, tend to build overly specific splitting rules. This forces the algorithm to perfectly memorize the noise and outliers within the training data instead of discovering underlying mathematical patterns.

To visually and statistically isolate this behavior, an unconstrained baseline Decision Tree Regressor was trained on an 80/20 data partition. The diagnostic testing revealed a clear validation gap:

- **Initial Training RMSE:** 0.0000
- **Initial Testing RMSE:** 0.4930

The fact that the training error dropped down to absolute perfection (0.0000) while the test error spiked up to 0.4930 statistically confirms severe overfitting. This initial model memorized unique variations in the training pool rather than capturing broader pricing structures, establishing a critical need for structural regularization.

3. Methodology & Validation Framework

To mitigate overfitting liabilities and establish an un-biased performance index, a professional validation pipeline was integrated:

- **Standardized Resampling (5-Fold Cross-Validation):** Rather than evaluating model stability using a single, arbitrary train-test split, the scaled

feature domain was parsed into 5 distinct, equal-sized subsets. The pipeline iteratively trained the estimator on 4 folds while validating on the remaining fold. This guarantees that every record is utilized for both training and validation, generating a realistic estimation of real-world error. Our unconstrained model yielded an aggregate 5-Fold Cross-Validated RMSE of 0.8962.

- **Automated Grid Optimization (GridSearchCV):** To control tree complexity, a hyperparameter grid sweep was executed over structural constraints:
 - `max_depth`: Evaluated bounds [3, 5, 7, 10] to strictly cap how deep the tree can grow.
 - `min_samples_split`: Evaluated minimum sample node requirements of [2, 5, 10] to smooth leaf transitions and ignore localized noise.

The optimization sweep systematically performed $4 \times 3 = 12$ distinct experiments across 5 folds each, computing a grand total of 60 individual training loops to locate the absolute best parameter mix.

4. Results & Hyperparameter Discovery

The cross-validated grid search successfully isolated the mathematically optimal structural hyperparameters for the California Housing Dataset:

- **Optimal `max_depth`: 10**
- **Optimal `min_samples_split`: 10**

Extracting this optimized model structure effectively constrained unmanageable branching. The final benchmarking summary table highlights the performance of the tuned tree relative to your previous baseline architectures:

Model Architecture	Valuation Framework	Root Mean Squared Error (RMSE)	R2 Score
Linear Regression (Baseline)	Single Train-Test Split	0.55590	0.575800
Ridge Regression	Single Train-Test Split	0.55580	0.575800
Tuned Decision Tree	5-Fold Cross-Validation	0.41658	0.682099

5. Model Selection Justification & Engineering Conclusions

Based on the empirical evidence gathered across the multi-stage pipeline, the **Tuned Decision Tree Regressor** is chosen as the final operational model for deployment.

Technical Justification Summary:

1. **Overfitting Control:** By locking structural thresholds to a maximum depth of 10 and requiring a split threshold of 10 samples, the model's overfitting liabilities were entirely mitigated. The final testing error (0.4166) tracks safely alongside cross-validated expectations, proving generalization stability.
2. **Explanatory Dominance:** The tuned model achieved a significantly lower RMSE of **0.41658** and a superior R^2 score of **0.682099** relative to both Linear and regularized Ridge alternatives. This proves that mapping localized multi-variable splits captures complex, non-linear pricing relationships (such as geographic clustering or population impact) far more effectively than linear weight functions.
3. **Statistical Trust:** Because this model's selection is backed by cross-validation metrics rather than a singular data split, its performance estimation is insulated against data anomalies. It represents a balanced trade-off between architectural complexity and real-world predictive reliability.